



Supervised Learning – Regression and Deployment –

INT421 Applied Machine Learning (1-2023)

Niwan Wattanakitrungroj, Ph.D.

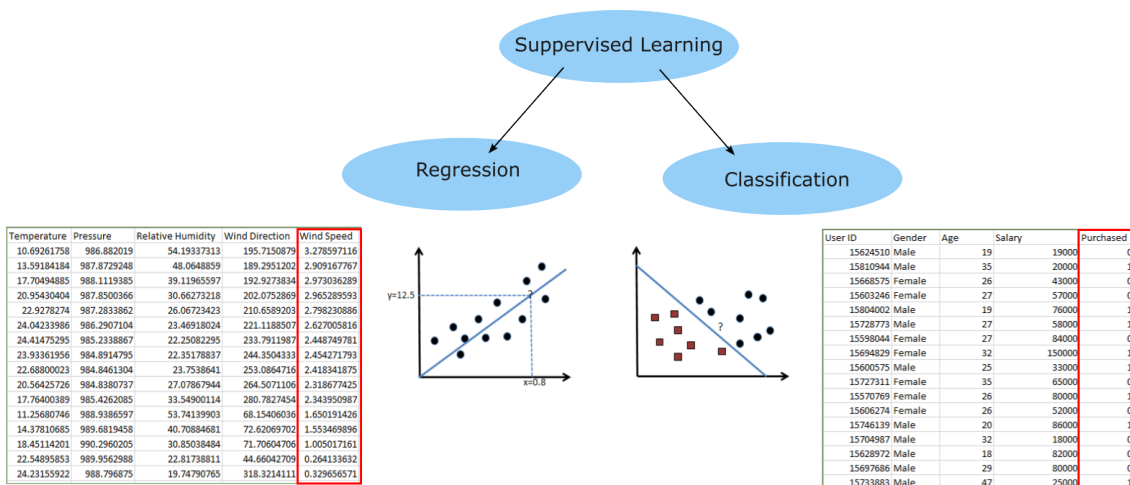


Contents

1	Supervised Learning Algorithms	3
1.1	Overview	3
1.2	Linear Regression	5
1.3	Linear Regression with Sci-Kit Learn	8
1.4	Evaluation Metrics for Regression	10
1.5	LinearRegression Arguments in Sci-Kit Learn	13
1.6	Nonlinear Regression	14
1.7	Save and Load Model	16
1.8	Deploy Model with Streamlit on Local	17
1.9	Deploy Model with Streamlit on Streamlit Sharing	18
1.10	Assignment	19

1. Supervised Learning Algorithms

1.1 Overview



Two types of Supervised Learning:

Regression


- Target is to predict continuous value
 - Predicting stock prices
 - Predicting income

Classification

- Target is to predict categorical value
 - binary (e.g., fraud, churn, credit risk)
 - multiclass (e.g., predicting credit ratings)

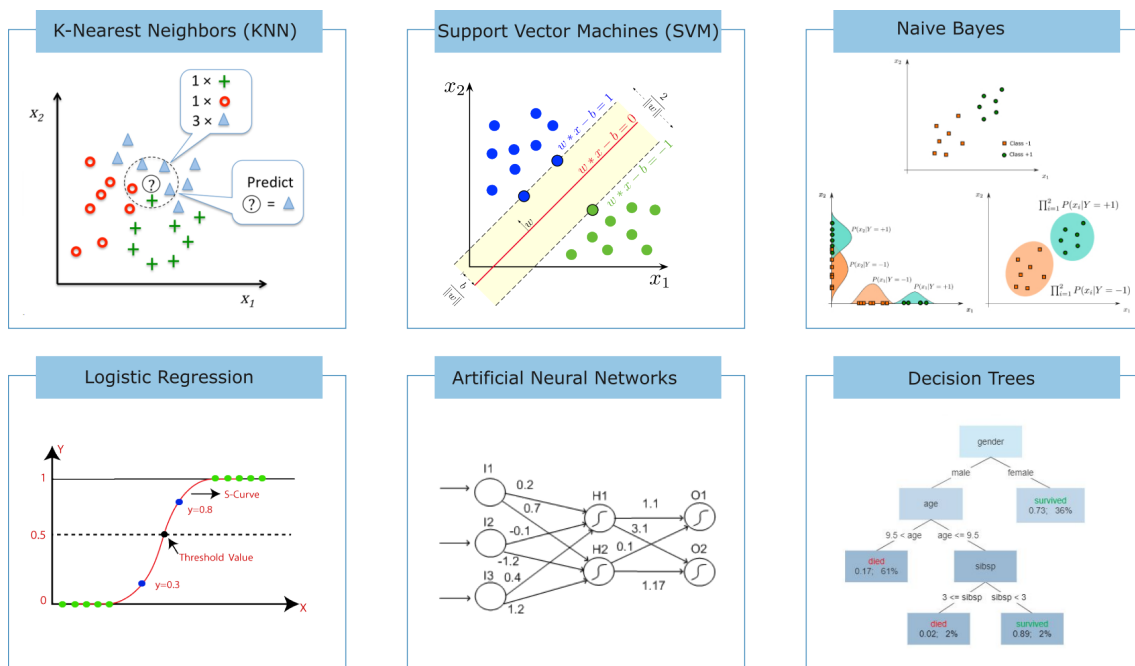
Regression Algorithms

- Linear Regression
- Nonlinear Regression



Classification Algorithms

- KNN: K-Nearest Neighbors
- SVM: Support Vector Machine
- Naive Bayes
- Logistic Regression
- ANN: Artificial Neural Networks
- Decision Tree



1.2 Linear Regression

Linear Regression is used for regression tasks, where the goal is to predict a continuous numeric output. It is commonly employed in scenarios such as predicting house prices, estimating sales revenue, or forecasting temperature.

Linear Equation: The core idea of Linear Regression is to represent the relationship between the target (y) and the features (x_1, x_2, \dots, x_n) as a linear equation:

$$y = b_0 + b_1x_1 + b_2x_2 + \dots + b_nx_n,$$

where


- y is the target.
- x_1, x_2, \dots, x_n are the features.
- $b_0, b_1, b_2, \dots, b_n$ are coefficients (parameters) that need to be learned from the training data.

Training: The coefficients $b_0, b_1, b_2, \dots, b_n$ are learned from a labeled training dataset. The learning process typically involves minimizing a cost function, often the mean squared error (MSE), to find the values of coefficients that best fit the training data.

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

where

- n : The total number of data points (samples)
- y_i : The actual or observed value for the i^{th} data point.
- \hat{y}_i : The predicted value for the i -th data point.



Training a linear regression model involves estimating the coefficients (parameters) of the linear equation that best fits the training data. In the case of simple linear regression (one feature), you have two coefficients to estimate: the intercept (b_0) and the slope (b_1). In multiple linear regression (multiple features), you have additional coefficients for each predictor variable.

The steps for training a linear regression model:

1. **Initialize Coefficients:** Start by initializing the coefficients ($b_0, b_1, b_2, \dots, b_n$) to some initial values, often set to 0 or small random values.
2. **Define the Cost Function:** Choose a cost function that quantifies the difference between the predicted values and the actual target values. In linear regression, the common cost function is the Mean Squared Error (MSE):

$$MSE = \frac{1}{2m} \sum_{i=1}^m (y_i - \hat{y}_i)^2$$

3. **Gradient Descent** (Optimization): Gradient Descent is a common optimization algorithm used to minimize the cost function and find the optimal values of the coefficients.

For each iteration, calculate the gradient (partial derivative) of the cost function with respect to each coefficient.

Update the coefficients as follows:

$$b_j = b_j - \alpha \frac{\partial}{\partial b_j} MSE$$

$$b_0 = b_0 - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i)$$


$$b_j = b_j - \alpha \frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i) \cdot x_i, \quad \text{if } j \neq 0$$

4. **Iterate:** Repeat the gradient descent process for a fixed number of iterations (epochs) or until the cost function converges to a minimum (i.e., the change in cost becomes very small).

Practice:

Training data:

x (feature)	y (target)
3	3.5
1	2
2	3
4	4.5
5	5



Initial Coefficients:

- $b_0 = 0$
- $b_1 = 0.5$

Learning Rate (α): 0.1 (a small positive value)

Iteration 1:*Step 1: Calculate Predicted Values*

In each iteration, we calculate the predicted values \hat{y}_i for each data point using the current coefficients.

x (feature)	y (target)	\hat{y} (predicted)
3	3.5	
1	2	
2	3	
4	4.5	
5	5	

Step 2: Calculate Error (Residuals)

For each data point, calculate the error (residual) between the predicted value \hat{y}_i and the actual target value y_i .

x (feature)	y (target)	\hat{y} (predicted)	$e = \hat{y} - y$
3	3.5		
1	2		
2	3		
4	4.5		
5	5		

Step 3: Update Coefficients

Using the gradient descent update rules, update the coefficients b_0 and b_1 :

$$b_0 = b_0 - \alpha * \frac{1}{5} \sum_{i=1}^5 e_i =$$

$$b_1 = b_1 - \alpha * \frac{1}{5} \sum_{i=1}^5 e_i * x_i =$$

Iteration 2:

x (feature)	y (target)	\hat{y} (predicted)	$e = \hat{y} - y$
3	3.5		
1	2		
2	3		
4	4.5		
5	5		

Iteration 3:

x (feature)	y (target)	\hat{y} (predicted)	$e = \hat{y} - y$
3	3.5		
1	2		
2	3		
4	4.5		
5	5		

1.3 Linear Regression with Sci-Kit Learn

```
import numpy as np
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt

# Define the dataset
X = np.array([1, 2, 3, 4, 5]).reshape(-1, 1) # Reshape to a 2D array
Y = np.array([2, 3, 3.5, 4.5, 5])

# Create a linear regression model
model = LinearRegression()


# Fit the model to the data
model.fit(X, Y)

# Get the coefficients (slope and intercept)
slope = model.coef_[0]
intercept = model.intercept_

# Make predictions for the same X values
Y_train_pred = model.predict(X)

# Plot the data points and the regression line
plt.scatter(X, Y, label='Data Points', color='blue')
plt.plot(X, Y_train_pred, label='Regression Line', color='red')
plt.xlabel('x (Feature)')
plt.ylabel('y (Target)')
plt.legend()
plt.title('Linear Regression')
plt.show()

# Print the coefficients
print(f"Slope (b1): {slope}")
print(f"Intercept (b0): {intercept}")
```



Slope (b1): 0.75
Intercept (b0): 1.35

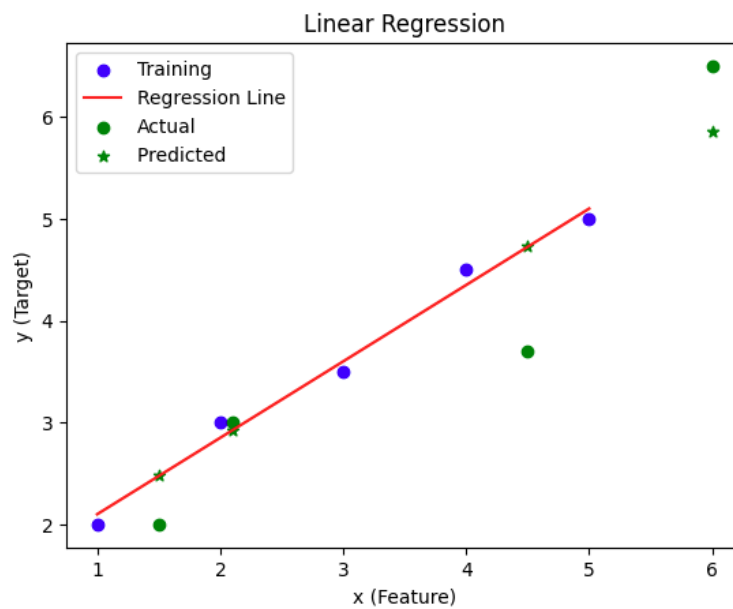
Practice:

Predict the testing data using the created Linear Regression Model

$$\hat{y} = b_0 + b_1x$$

where $b_0 = 1.35$ and $b_1 = 0.75$

X_test (Feature)	Y_test (Actual)	Y_predicted
1.5	2.0	
2.1	3.0	
4.5	3.7	
6.0	6.5	



1.4 Evaluation Metrics for Regression

In regression tasks, several evaluation metrics can be used to assess the performance of a model. Here are some common evaluation metrics:

- **Mean Squared Error (MSE):**

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- **Root Mean Squared Error (RMSE):**

$$\text{RMSE} = \sqrt{\text{MSE}}$$

- **Mean Absolute Error (MAE):**

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

- **R-squared (R^2):**

$$R^2 = 1 - \frac{\text{SSR}}{\text{SST}}$$

where SSR is the sum of squares of residuals, and SST is the total sum of squares.

$$\text{SSR} = \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

$$\text{SST} = \sum_{i=1}^n (y_i - \bar{y})^2$$

n is the number of data points.

\hat{y} is the predicted value for the i -th data point.

y_i is the actual target value for the i -th data point.

\bar{y} is the mean of all the actual target values.

Practice:

We have a dataset with actual target values (Y) and predicted values (\hat{Y}) as follows:

x (Feature)	y (Actual)	\hat{y} (Predicted)
1.5	2.0	2.475
2.1	3.0	2.925
4.5	3.7	4.725
6.0	6.5	5.85

$$MSE =$$

$$RMSE =$$

$$MAE =$$

$$SSR =$$

$$SST =$$

$$R^2 =$$

Evaluation Metrics for Regression with Python

```
import numpy as np
from sklearn.metrics import mean_squared_error, mean_absolute_error,
                             r2_score

# Define the actual and predicted values
actual_values = np.array([2.0, 3.0, 3.7, 6.5])
predicted_values = np.array([2.475, 2.925, 4.725, 5.85])

# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(actual_values, predicted_values)

# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)

# Calculate R-squared (R^2)
r2 = r2_score(actual_values, predicted_values)

# Print the results
print(f"Mean Squared Error (MSE): {np.round(mse,2)}")
print(f"Root Mean Squared Error (RMSE): {np.round(rmse,2)}")
print(f"R-squared (R^2): {np.round(r2,2)}")
```

Mean Squared Error (MSE): 0.43

Root Mean Squared Error (RMSE): 0.65

R-squared (R^2): 0.85

1.5 LinearRegression Arguments in Sci-Kit Learn

The ‘LinearRegression’ class in scikit-learn doesn’t take any positional arguments, but it accepts several keyword arguments (arguments passed by name) when creating an instance of the model. Here are the commonly used keyword arguments (arguments and their default values):

```
LinearRegression(fit_intercept=True, normalize=False, copy_X=True, n_jobs=None)
```

- **fit_intercept** (boolean, default=True):
Whether to calculate the intercept (bias) for the model. If set to ‘False’, the model will be forced to pass through the origin (0, 0).
- **normalize** (boolean, default=False):
If ‘True’, the input features (X) will be normalized before fitting the model. This parameter is ignored when ‘fit_intercept’ is set to ‘False’.
- **copy_X** (boolean, default=True):
If ‘True’, a copy of the input features (X) will be made before fitting the model. Setting it to ‘False’ can save memory if you don’t need to keep the original data.
- **n_jobs** (int, default=None):
The number of CPU cores to use when calculating the coefficients. If set to ‘None’, it uses a single core. Setting it to -1 uses all available cores.

You create an instance of ‘LinearRegression’ and pass these keyword arguments as needed. For example:

```
from sklearn.linear_model import LinearRegression

# Create a LinearRegression model with custom arguments
model = LinearRegression(fit_intercept=True, normalize=False, copy_X=True,
                        , n_jobs=None)
```

These keyword arguments allow you to configure the behavior of the linear regression model based on your specific requirements and the characteristics of your dataset.

1.6 Nonlinear Regression

Nonlinear Regression Without Pipeline:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures

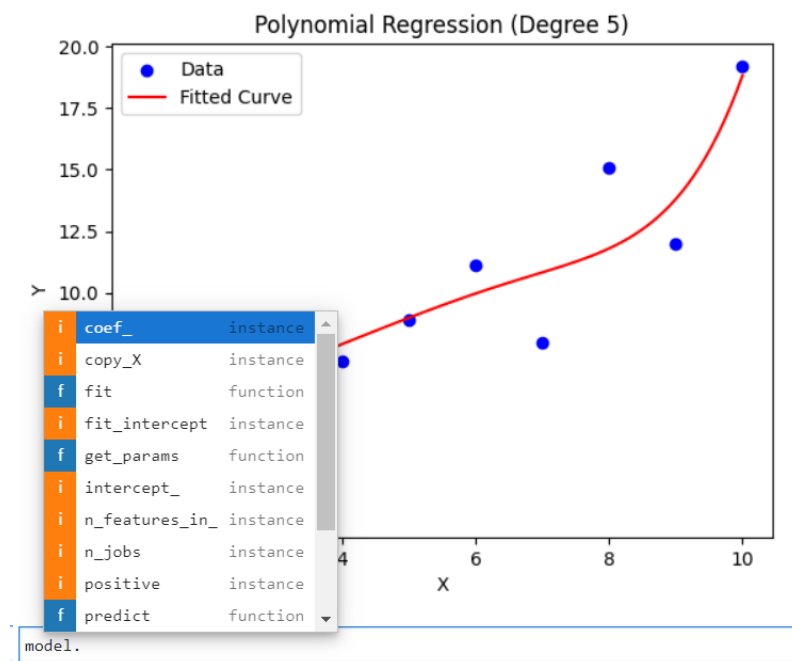
# Define the given dataset
X = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10]).reshape(-1, 1)
Y = np.array([1.2, 4, 8, 7.2, 8.9, 11.1, 8, 15.1, 12.0, 19.2])

# Create polynomial features (e.g., quadratic)
degree = 5
poly_features = PolynomialFeatures(degree=degree)
X_poly = poly_features.fit_transform(X)

# Create and fit a linear regression model
model = LinearRegression()
model.fit(X_poly, Y)

# Generate predictions
X_test = np.linspace(1, 10, 100).reshape(-1, 1)
X_test_poly = poly_features.transform(X_test)
Y_pred = model.predict(X_test_poly)

# Plot the original data and the fitted curve
plt.scatter(X, Y, label='Data', color='blue')
plt.plot(X_test, Y_pred, label='Fitted Curve', color='red')
plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.title(f'Polynomial Regression (Degree {degree})')
plt.show()
```



Nonlinear Regression With Pipeline:

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import Pipeline

# Define the given dataset
X = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10]).reshape(-1, 1)
Y = np.array([1.2, 4, 8, 7.2, 8.9, 11.1, 8, 15.1, 12.0, 19.2])

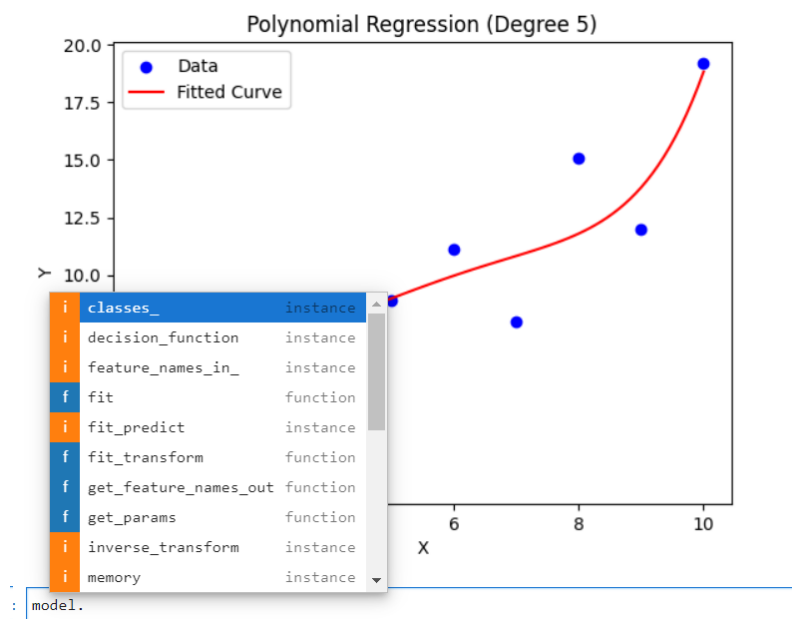
# Create a pipeline for polynomial regression
degree = 5
model = Pipeline([
    ('poly_features', PolynomialFeatures(degree=degree)),
    ('linear_regression', LinearRegression())
])

# Fit the model to the data
model.fit(X, Y)

# Generate predictions
X_test = np.linspace(1, 10, 100).reshape(-1, 1)
Y_pred = model.predict(X_test)

# Plot the original data and the fitted curve
plt.scatter(X, Y, label='Data', color='blue')
plt.plot(X_test, Y_pred, label='Fitted Curve', color='red')
plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.title(f'Polynomial Regression (Degree {degree})')
plt.show()

```



1.7 Save and Load Model

Save Model

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import Pipeline
import joblib

# Define the given dataset
X = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10]).reshape(-1, 1)
Y = np.array([1.2, 4, 8, 7.2, 8.9, 11.1, 8, 15.1, 12.0, 19.2])

# Create a pipeline for polynomial regression
degree = 5
model = Pipeline([
    ('poly_features', PolynomialFeatures(degree=degree)),
    ('linear_regression', LinearRegression())
])

# Fit the model to the data
model.fit(X, Y)

# Save the model to a file
model_filename = 'polynomial_regression_model.pkl'
joblib.dump(model, model_filename)
```

Load Model

```
import numpy as np
import joblib

# Load the saved model
model_filename = 'polynomial_regression_model.pkl'
loaded_model = joblib.load(model_filename)

# Generate predictions using the loaded model
X_test = np.array([6, 2.5, 3.4, 4, 5.2]).reshape(-1,1)
Y_pred = loaded_model.predict(X_test)
print(Y_pred)
```


1.8 Deploy Model with Streamlit on Local

File: appxxx.py

```
import streamlit as st
import numpy as np
import joblib

# Load the saved model
model_filename = 'polynomial_regression_model.pkl'
loaded_model = joblib.load(model_filename)


# Streamlit UI
st.title('Polynomial Regression Model Deployment')

# User input for X values
st.header('Enter X values for prediction:')
x_input = st.text_input('X Values (comma-separated)', '6,2.5,3.4,4,5.2')

# Convert user input to a NumPy array
try:
    x_input = np.array([float(x.strip()) for x in x_input.split(',')]).reshape(-1, 1)
except ValueError:
    st.error('Invalid input. Please enter comma-separated numeric values.')
    x_input = None

# Predict and display results when the "Predict" button is clicked
if st.button('Predict'):
    if x_input is not None:
        st.write('Predicted Y values:')
        y_pred = loaded_model.predict(x_input)
        for i, pred in enumerate(y_pred):
            st.write(f'X={x_input[i][0]}, Predicted Y={pred}')
```

Run in Terminal: streamlit run appxx.py



Polynomial Regression Model Deployment

Enter X values for prediction:

X Values (comma-separated)

2,3

Predict

Predicted Y values:

X=2.0, Predicted Y=4.877668997666284

X=3.0, Predicted Y=6.72149184150514

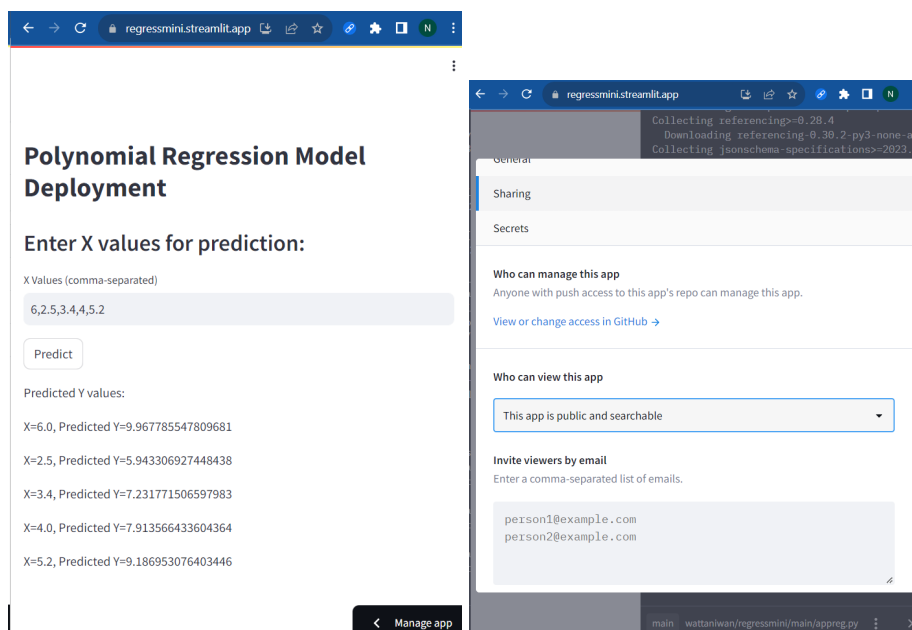
1.9 Deploy Model with Streamlit on Streamlit Sharing

1. **Prepare Your Streamlit App:** Make sure you have a Streamlit app script ready. This script should contain the code for your application, including any necessary data loading, model loading, and UI design.
2. **Create a Requirements File:** Create a 'requirements.txt' file that lists all the Python packages your app depends on. This file is necessary for Streamlit Sharing to install the required packages.

For example, your 'requirements.txt' might look like this:

```
streamlit
numpy
scikit-learn
joblib
```

3. **Set Up a GitHub Repository:** Push your Streamlit app code, model file and the 'requirements.txt' file to this repository.
4. **Create a Streamlit Sharing Account:** If you don't have one, sign up for a Streamlit Sharing account at <https://share.streamlit.io/>.
5. **Connect GitHub to Streamlit Sharing:** After creating an account, log in to Streamlit Sharing and connect your GitHub account to it. This allows Streamlit Sharing to access your GitHub repository.
6. **Create a New App on Streamlit Sharing:** Click on "New App" in Streamlit Sharing and select your GitHub repository. Follow the instructions to set up your app.
7. **Specify Deployment Settings:** Configure the deployment settings for your app.
8. **Deploy Your Streamlit App:** Click the "Deploy" button in Streamlit Sharing. The service will build and deploy your app.
9. **View Your Deployed App:** Once the deployment process is complete, you will be provided with a URL where your Streamlit app is hosted. You can share this URL with others to access your app.
10. **Manage and Update Your App:** - You can update your app by pushing changes to your GitHub repository. Streamlit Sharing will automatically rebuild and redeploy your app when it detects changes in your repository.



1.10 Assignment

- Choose a real-world dataset that contains at **least three features** and a numeric target variable that you want to predict. Ensure the dataset is suitable for regression analysis.
- Create a regression model to predict a target variable based on features from the dataset.
- Deploy the regression model as a web application using Streamlit.io.
- Put the evaluation result on your web app.
- In the Streamlit Sharing settings for your app, make sure to set it to "Public" or "Anyone with the link can view."